

Our exploratory data analysis examined the housing market in Ames, Iowa, using the house price [SalePrice] as the dependent variable of interest. The initial dataset includes 1460 data records and 81 fields. The test data has a mean sale price of \$180,921.20 and a standard deviation of \$79442.50. We cleaned the data by removing fields with null values and trimmed the dataset using an isolation forest algorithm to detect and remove 5% of the anomalies in [SalePrice]. We also removed all non-numeric features. The cleaned test data has 1,387 data records, 37 fields, a median sale price of \$173,354.52, and a standard deviation of \$59,601.46. (Previously, we removed outliers more than two standard deviations away and ended with 1,459 data records, a median sale price of \$165,228.32, and a standard deviation of \$58,943.80.) The isolation forest algorithm improved our models' accuracy but decreased the correlation coefficient for all predictor variables.

We engineered four new features to use as predictor variables. A [baths] feature took [FullBath] plus half of [HalfBath]; and a [newnes_value] feature took the [YearBuilt] plus [YearRemodAdd] divided by [YearBuilt], giving older, remodeled homes a modernity adjustment. We also continued to use a total square foot variable [tot_sq] that measures the houses' total square footage (including basements) and a quality space [qual_space] variable that multiplies [total_sq] by overall quality [OverallQual]. See Figure 1 in the index for correlation coefficients of some variables used in our models.

Following our exploratory data analysis, we did feature selection. Our linear regression models generally used int64 features from the original dataset and features that we engineered. Our Lasso, Ridge, and ElasticNet regression used the 12 most important numeric features as determined by recursive feature elimination. We prepared all of our models with a cross-validation design. We divided our training data set into two groups: a training data set of 932 (80%) and a testing data set of 23 (20%). We tested all of our models and chose the linear

regression model that best predicted the “test” training data and all three machine learning models to use with the “real” test data.

Using the k-fold cross-validation method, we built four regression models using the training data to predict the [SalePrice] for each house: Linear (OLS), Lasso, Ridge, and ElasticNet regression models. Linear regression Model One has the best fit out of the four linear regression models since it has the highest mean r-squared score of 0.70 and the lowest standard deviation of 0.10 (slightly better than the Ridge Regression). Our other regression models all had increased r-squared scores and larger standard deviations. Our Ridge Regression is our best model did best with the test data and had a mean r-squared score of 0.78 and a standard deviation of 0.12. The Lasso Regression had a mean r-squared score of 0.76 and a standard deviation of 0.18. Finally, the initial ElasticNet Regression model has a mean r-squared score of 0.47 and a standard deviation of 0.11. We tuned our hyperparameter for this model using an $\alpha = 0.10$ and $\text{l1_ratio} = 0.75$. This resulted in increasing the mean r-squared score to 0.76 and the standard deviation to 0.18.

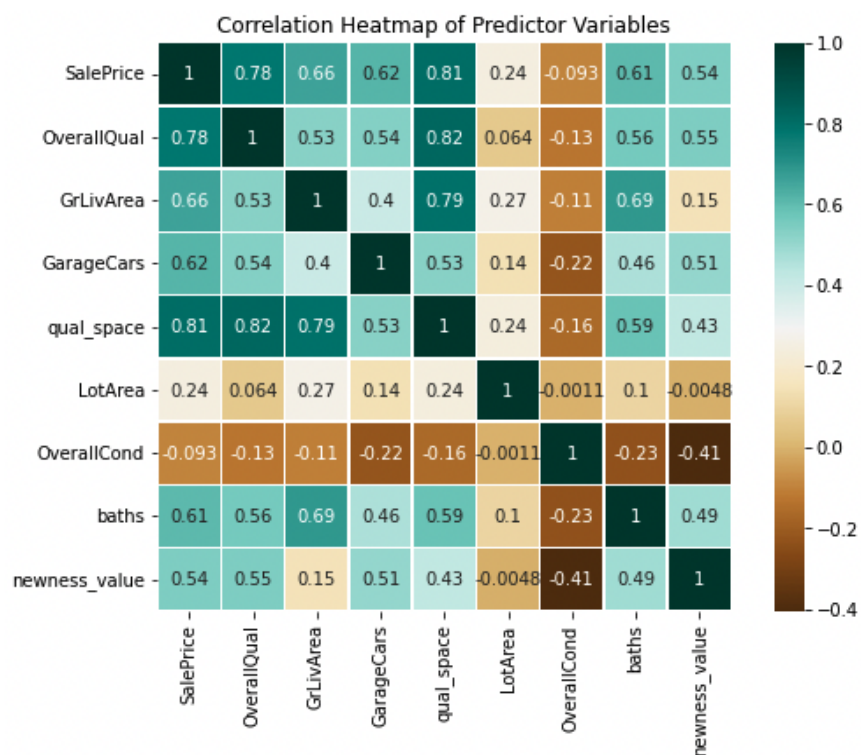
Our ElasticNet Regression model was determined to be the best predictor based on its Kaggle score of 0.18135. Our three machine learning regression models received similar r-squares scores and standard deviations, so it's not too much of a surprise that the ElasticNet Regression model did better with the actual test data, while the Ridge Regression model did best with the training “test” data. Our other regression models ranked in the following order: Ridge, 0.18569; Lasso, 0.18589; Linear, 0.20912. Our previous linear model received a slightly better Kaggle score of 0.20317, but all of our machine learning models beat our previous Kaggle score.

Index:

Kaggle user name for upload: ZachWat

✓	group_5_elasticnet_m3 (1).csv Submitted by ZachWat · Submitted just now	Score: 0.18135
✓	group_5_lasso_m3 (1).csv Submitted by ZachWat · Submitted just now	Score: 0.18589
✓	group_5_ridge_m3 (1).csv Submitted by ZachWat · Submitted just now	Score: 0.18569
✓	group_5_linear_m3 (1).csv Submitted by ZachWat · Submitted 2 minutes ago	Score: 0.20912

Figure 1:



Intro

Links

<https://canvas.northwestern.edu/courses/167719/assignments/1078600>

<https://www.kaggle.com/c/house-prices-advanced-regression-techniques>

Modules

```
In [1]: #For data manipulation and visualization
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

#For Isolation Forest from sklearn
from sklearn.ensemble import IsolationForest
from enum import auto

#Models from sklearn (Linear, Lasso, and Ridge)
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge

#ElasticNet from sklearn
from sklearn.linear_model import ElasticNet
from sklearn.linear_model import ElasticNetCV
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RepeatedKFold
from numpy import arange

#Other from sklearn
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error

#Other
from math import sqrt
import warnings
```

Import Files

```
In [2]: #Import train.csv and test.csv from the Kaggle page Linked above
#from google.colab import files
#uploaded = files.upload()
```

```
In [3]: df = pd.read_csv("train.csv")
```

EDA

Intro Stats

In [4]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Id                    1460 non-null   int64
 1   MSSubClass            1460 non-null   int64
 2   MSZoning              1460 non-null   object
 3   LotFrontage          1201 non-null   float64
 4   LotArea               1460 non-null   int64
 5   Street               1460 non-null   object
 6   Alley                91 non-null     object
 7   LotShape              1460 non-null   object
 8   LandContour          1460 non-null   object
 9   Utilities             1460 non-null   object
10   LotConfig            1460 non-null   object
11   LandSlope             1460 non-null   object
12   Neighborhood          1460 non-null   object
13   Condition1           1460 non-null   object
14   Condition2           1460 non-null   object
15   BldgType              1460 non-null   object
16   HouseStyle            1460 non-null   object
17   OverallQual           1460 non-null   int64
18   OverallCond           1460 non-null   int64
19   YearBuilt             1460 non-null   int64
20   YearRemodAdd         1460 non-null   int64
21   RoofStyle            1460 non-null   object
22   RoofMatl             1460 non-null   object
23   Exterior1st          1460 non-null   object
24   Exterior2nd          1460 non-null   object
25   MasVnrType           1452 non-null   object
26   MasVnrArea           1452 non-null   float64
27   ExterQual            1460 non-null   object
28   ExterCond            1460 non-null   object
29   Foundation           1460 non-null   object
30   BsmtQual             1423 non-null   object
31   BsmtCond             1423 non-null   object
32   BsmtExposure         1422 non-null   object
33   BsmtFinType1         1423 non-null   object
34   BsmtFinSF1           1460 non-null   int64
35   BsmtFinType2         1422 non-null   object
36   BsmtFinSF2           1460 non-null   int64
37   BsmtUnfSF            1460 non-null   int64
38   TotalBsmtSF          1460 non-null   int64
39   Heating              1460 non-null   object
40   HeatingQC            1460 non-null   object
41   CentralAir           1460 non-null   object
42   Electrical           1459 non-null   object
43   1stFlrSF             1460 non-null   int64
44   2ndFlrSF             1460 non-null   int64
45   LowQualFinSF         1460 non-null   int64
46   GrLivArea            1460 non-null   int64
47   BsmtFullBath         1460 non-null   int64
48   BsmtHalfBath         1460 non-null   int64
49   FullBath             1460 non-null   int64
50   HalfBath             1460 non-null   int64
```

```

51 BedroomAbvGr 1460 non-null int64
52 KitchenAbvGr 1460 non-null int64
53 KitchenQual 1460 non-null object
54 TotRmsAbvGrd 1460 non-null int64
55 Functional 1460 non-null object
56 Fireplaces 1460 non-null int64
57 FireplaceQu 770 non-null object
58 GarageType 1379 non-null object
59 GarageYrBlt 1379 non-null float64
60 GarageFinish 1379 non-null object
61 GarageCars 1460 non-null int64
62 GarageArea 1460 non-null int64
63 GarageQual 1379 non-null object
64 GarageCond 1379 non-null object
65 PavedDrive 1460 non-null object
66 WoodDeckSF 1460 non-null int64
67 OpenPorchSF 1460 non-null int64
68 EnclosedPorch 1460 non-null int64
69 3SsnPorch 1460 non-null int64
70 ScreenPorch 1460 non-null int64
71 PoolArea 1460 non-null int64
72 PoolQC 7 non-null object
73 Fence 281 non-null object
74 MiscFeature 54 non-null object
75 MiscVal 1460 non-null int64
76 MoSold 1460 non-null int64
77 YrSold 1460 non-null int64
78 SaleType 1460 non-null object
79 SaleCondition 1460 non-null object
80 SalePrice 1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

In [5]: `df.head()`

Out[5]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	.
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	

5 rows × 81 columns



Descriptive statistics and visualizations to help understand the marginal distribution of the dependent variable.

In [6]: `df["SalePrice"].describe()`

Out[6]:

count	1460.000000
mean	180921.195890
std	79442.502883
min	34900.000000
25%	129975.000000
50%	163000.000000

```


75%      214000.000000
max      755000.000000
Name: SalePrice, dtype: float64

```

EDA SalePrice Graphs


```
In [7]: sns.histplot(x="SalePrice", data=df)
```

```
Out[7]: <AxesSubplot:xlabel='SalePrice', ylabel='Count'>
```



```
In [8]: sns.boxplot(x="SalePrice", data=df)
```

```
Out[8]: <AxesSubplot:xlabel='SalePrice'>
```



Investigate Missing Data and Outliers

Missing Data:

```
In [9]: #Toggle this to see results below
pd.set_option('max_rows', 20)
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: Id          0
      MSSubClass    0
      MSZoning      0
      LotFrontage   259
      LotArea       0
      ...
      MoSold        0
      YrSold        0
      SaleType      0
      SaleCondition 0
      SalePrice     0
      Length: 81, dtype: int64
```

The following categories have null values:

- LotFrontage
- Alley
- MasVnrType
- MasVnrArea
- BsmtQual
- BsmtCond
- BsmtExposure
- BsmtFinType1
- BsmtFinType2
- Electrical
- FireplaceQu
- GarageType
- GarageYrBlt
- GarageFinish
- GarageQual
- GarageCond
- PoolQC
- Fence
- MiscFeature

We're not concerned about most of these columns having null values. It makes sense that some of the data would be missing for each (if a house doesn't have a pool, for example). We're going to drop all of the columns that have null values with the exception of "Electrical." For 'Electrical,' we'll remove the row with the null value.

```
In [11]: col_to_drop = ['LotFrontage', 'Alley', 'MasVnrType', 'MasVnrArea',
                        'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', '
                        'GarageFinish', 'GarageQual', 'GarageCond', 'PoolQC', 'Fence', 'MiscFeat

df2a = df.drop(columns=col_to_drop, inplace=False)

col_to_drop_2 = ['MSZoning', 'Street', 'LotShape', 'LandContour', 'Utilities', 'LotConfig
                'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'R
                'RoofMat1', 'Exterior1st', 'Exterior2nd', 'ExterQual', 'ExterCond', 'Fo
                'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual', 'Func
                'PavedDrive', 'SaleType', 'SaleCondition']

df2b = df2a.drop(columns= col_to_drop_2, inplace=False)
```



```
df3 = df2b.dropna()
df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1460 entries, 0 to 1459
Data columns (total 35 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                     1460 non-null   int64
1   MSSubClass             1460 non-null   int64
2   LotArea                1460 non-null   int64
3   OverallQual            1460 non-null   int64
4   OverallCond            1460 non-null   int64
5   YearBuilt              1460 non-null   int64
6   YearRemodAdd           1460 non-null   int64
7   BsmtFinSF1             1460 non-null   int64
8   BsmtFinSF2             1460 non-null   int64
9   BsmtUnfSF              1460 non-null   int64
10  TotalBsmtSF            1460 non-null   int64
11  1stFlrSF               1460 non-null   int64
12  2ndFlrSF               1460 non-null   int64
13  LowQualFinSF           1460 non-null   int64
14  GrLivArea              1460 non-null   int64
15  BsmtFullBath           1460 non-null   int64
16  BsmtHalfBath           1460 non-null   int64
17  FullBath               1460 non-null   int64
18  HalfBath               1460 non-null   int64
19  BedroomAbvGr           1460 non-null   int64
20  KitchenAbvGr           1460 non-null   int64
21  TotRmsAbvGrd           1460 non-null   int64
22  Fireplaces             1460 non-null   int64
23  GarageCars             1460 non-null   int64
24  GarageArea             1460 non-null   int64
25  WoodDeckSF             1460 non-null   int64
26  OpenPorchSF            1460 non-null   int64
27  EnclosedPorch          1460 non-null   int64
28  3SsnPorch              1460 non-null   int64
29  ScreenPorch            1460 non-null   int64
30  PoolArea               1460 non-null   int64
31  MiscVal                1460 non-null   int64
32  MoSold                 1460 non-null   int64
33  YrSold                 1460 non-null   int64
34  SalePrice              1460 non-null   int64
dtypes: int64(35)
memory usage: 410.6 KB
```

Removing Anomalies

```
In [12]: #Isolation Forest Removes anomalies
model=IsolationForest(n_estimators=100, contamination=float(.05), random_state=42)

model.fit(df3[['SalePrice']])

print(model.get_params())

{'bootstrap': False, 'contamination': 0.05, 'max_features': 1.0, 'max_samples': 'auto',
'n_estimators': 100, 'n_jobs': None, 'random_state': 42, 'verbose': 0, 'warm_start': False}
```

```
In [13]: df3['scores'] = model.decision_function(df3[['SalePrice']])
```

```
df3['anomaly_score'] = model.predict(df3[['SalePrice']])

df4 = df3[df3['anomaly_score']!= -1]

df4.shape
```

Out[13]: (1387, 37)


EDA SalePrice Stats/Graphs (Cleaned Data)

```
In [14]: df4["SalePrice"].describe()
```

```
Out[14]: count      1387.000000
mean      173354.517664
std       59601.459217
min       72500.000000
25%      130000.000000
50%      160200.000000
75%      206600.000000
max      361919.000000
Name: SalePrice, dtype: float64
```


```
In [15]: sns.histplot(x="SalePrice", data=df4)
```

Out[15]: <AxesSubplot:xlabel='SalePrice', ylabel='Count'>



```
In [16]: sns.boxplot(x="SalePrice", data=df4)
```

Out[16]: <AxesSubplot:xlabel='SalePrice'>



Feature Engineering

In [17]:

```
# sum 1st floor, 2nd floor, and basement square footage to get total square footage
sum_column = df4['1stFlrSF'] + df4['2ndFlrSF'] + df4['TotalBsmtSF']

# multiply total square footage by overall quality to generate new predictor variable q
mult_column = sum_column*df4['OverallQual']

#Sum the baths and the halfbaths
bath_column = df4['FullBath'] + 0.5*df4['HalfBath']

#Create a newness score (YearBuilt + (YearRemodAdd/YearBuilt))
newness_column = df4['YearBuilt'] + (df4['YearRemodAdd'] / df4['YearBuilt'])

# add new predictor variables to dataframe
df4['tot_sq'] = sum_column
df4['qual_space'] = mult_column
df4['baths'] = bath_column
df4['newness_value'] = newness_column
df4.head()
```

<ipython-input-17-d7942eafcf14>:16: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df4['tot_sq'] = sum_column
<ipython-input-17-d7942eafcf14>:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df4['qual_space'] = mult_column
<ipython-input-17-d7942eafcf14>:18: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df4['baths'] = bath_column
<ipython-input-17-d7942eafcf14>:19: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df4['newness_value'] = newness_column
```

```
Out[17]:
```

	Id	MSSubClass	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	BsmtFinSF1	BsmtFin
0	1	60	8450	7	5	2003	2003	706	
1	2	20	9600	6	8	1976	1976	978	
2	3	60	11250	7	5	2001	2002	486	
3	4	70	9550	7	5	1915	1970	216	
4	5	60	14260	8	5	2000	2000	655	

5 rows × 41 columns



Correlations

```
In [18]: # setting the columns to correlate


columns = ['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars', 'qual_space', 'LotArea']
df_corr = df4[columns]
# running the correlation
df_corr.corr()

# setting up the heatmap
corrmat = df_corr.corr()

# set the figure size
f, ax = plt.subplots(figsize=(9, 6))

# pass the data and set the parameters
sns.heatmap(corrmat, vmax=1, square=True, annot=True, cmap='BrBG', linewidths=.5 )
plt.title('Correlation Heatmap of Predictor Variables')

# images can be saved - default is .png
# https://matplotlib.org/3.1.1/api/_as_gen/matplotlib.pyplot.savefig.html
plt.savefig('Correlation Heatmap of New Predictor Variables')
```




Graphs

Scatterplots


```
In [19]: #sns.scatterplot(x='OverallQual', y='SalePrice', data=df4)
sns.lmplot(x='OverallQual', y='SalePrice', data=df4)
```

```
Out[19]: <seaborn.axisgrid.FacetGrid at 0x1990479ab80>
```




```
In [20]: #sns.scatterplot(x="GrLivArea", y="SalePrice", data=df4)
sns.lmplot(x="GrLivArea", y="SalePrice", data=df4)
```

```
Out[20]: <seaborn.axisgrid.FacetGrid at 0x19904600e20>
```




```
In [21]: #sns.scatterplot(x="GarageCars", y="SalePrice", data=df4)
sns.lmplot(x="GarageCars", y="SalePrice", data=df4)
```

```
Out[21]: <seaborn.axisgrid.FacetGrid at 0x19904600e50>
```




```
In [22]: #sns.scatterplot(x="GarageArea", y="SalePrice", data=df4)
sns.lmplot(x="GarageArea", y="SalePrice", data=df4)
```

```
Out[22]: <seaborn.axisgrid.FacetGrid at 0x199057b1ac0>
```




```
In [23]: #scatterplot total square ft vs. SalePrice
sns.lmplot(x="tot_sq", y="SalePrice", data=df4)
```

```
Out[23]: <seaborn.axisgrid.FacetGrid at 0x19903a61c10>
```




```
In [24]: #scatterplot qual_space vs. SalePrice
sns.lmplot(x="qual_space", y="SalePrice", data=df4)
```

Out[24]: <seaborn.axisgrid.FacetGrid at 0x1990587bb50>




```
In [25]: #scatterplot baths vs. SalePrice
sns.lmplot(x="baths", y="SalePrice", data=df4)
```

Out[25]: <seaborn.axisgrid.FacetGrid at 0x199058b64f0>



```
In [26]: # scatterplot of newness value vs SalePrice
sns.lmplot(x="newness_value", y="SalePrice", data=df4)
```


```
Out[26]: <seaborn.axisgrid.FacetGrid at 0x199058cf850>
```



Boxplots


```
In [27]: var = 'OverallQual'
data = pd.concat([df4['SalePrice'], df4[var]], axis=1)
sns.boxplot(x=var, y='SalePrice', data=data)
```

Out[27]: <AxesSubplot:xlabel='OverallQual', ylabel='SalePrice'>




```
In [28]: var = 'GarageCars'
data = pd.concat([df4['SalePrice'], df4[var]], axis=1)
sns.boxplot(x=var, y='SalePrice', data=data)
```

Out[28]: <AxesSubplot:xlabel='GarageCars', ylabel='SalePrice'>



```
In [29]: # side by side boxplot of baths vs. SalePrice
var = 'baths'
data = pd.concat([df4['SalePrice'], df4[var]], axis=1)
sns.boxplot(x=var, y='SalePrice', data=data)
```

Out[29]: <AxesSubplot:xlabel='baths', ylabel='SalePrice'>



Regression Models

Lasso, Ridge, and ElasticNet

Variable Selection

```
In [30]: z = df4.drop(columns = ['SalePrice', 'Id', 'scores', 'anomaly_score'])
y = df4['SalePrice']
```

```
In [31]: from sklearn.feature_selection import RFE
from sklearn.svm import SVR
estimator = SVR(kernel="linear")
selector = RFE(estimator, n_features_to_select=12, step=1)
selector = selector.fit(z, y)

df_ranks = selector.ranking_
df_ranks
```

```
Out[31]: array([[ 1, 24,  1,  1,  1,  1, 13, 18, 19,  6, 23, 16,  3,  1,  1, 21, 11,
                14,  1,  4, 25,  1,  9,  2,  8,  5, 10, 17,  1, 15, 26,  1,  7, 22,
                20, 12,  1])
```

```
In [32]: for i in range(z.shape[1]):
print('Column: %d, Selected %s, Rank: %.3f' % (i, selector.support_[i], selector.ra
```

```
Column: 0, Selected True, Rank: 1.000
Column: 1, Selected False, Rank: 24.000
Column: 2, Selected True, Rank: 1.000
Column: 3, Selected True, Rank: 1.000
Column: 4, Selected True, Rank: 1.000
Column: 5, Selected True, Rank: 1.000
Column: 6, Selected False, Rank: 13.000
Column: 7, Selected False, Rank: 18.000
Column: 8, Selected False, Rank: 19.000
Column: 9, Selected False, Rank: 6.000
Column: 10, Selected False, Rank: 23.000
Column: 11, Selected False, Rank: 16.000
Column: 12, Selected False, Rank: 3.000
```

```

Column: 13, Selected True, Rank: 1.000
Column: 14, Selected True, Rank: 1.000
Column: 15, Selected False, Rank: 21.000
Column: 16, Selected False, Rank: 11.000
Column: 17, Selected False, Rank: 14.000
Column: 18, Selected True, Rank: 1.000
Column: 19, Selected False, Rank: 4.000
Column: 20, Selected False, Rank: 25.000
Column: 21, Selected True, Rank: 1.000
Column: 22, Selected False, Rank: 9.000
Column: 23, Selected False, Rank: 2.000
Column: 24, Selected False, Rank: 8.000
Column: 25, Selected False, Rank: 5.000
Column: 26, Selected False, Rank: 10.000
Column: 27, Selected False, Rank: 17.000
Column: 28, Selected True, Rank: 1.000
Column: 29, Selected False, Rank: 15.000
Column: 30, Selected False, Rank: 26.000
Column: 31, Selected True, Rank: 1.000
Column: 32, Selected False, Rank: 7.000
Column: 33, Selected False, Rank: 22.000
Column: 34, Selected False, Rank: 20.000
Column: 35, Selected False, Rank: 12.000
Column: 36, Selected True, Rank: 1.000

```

```

In [33]: cols_to_keep = selector.support_

columns_to_remove = z.columns.values[np.logical_not(cols_to_keep)]
columns_to_remove

```

```

Out[33]: array(['LotArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
               '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'BsmtHalfBath', 'FullBath',
               'HalfBath', 'KitchenAbvGr', 'TotRmsAbvGrd', 'GarageCars',
               'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch',
               '3SsnPorch', 'PoolArea', 'MiscVal', 'YrSold', 'tot_sq',
               'qual_space', 'baths'], dtype=object)

```

```

In [34]: X = z.drop(columns=columns_to_remove)

```

```

In [35]: name_list = list(X.columns.values)

```

```

In [36]: # split data in to training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4

```

Lasso

```

In [37]: lasso_model = Lasso().fit(X_train, y_train)

```

```

C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 252181185368.77197, tolerance: 391652989.5874466
model = cd_fast.enet_coordinate_descent(

```

```

In [38]: # display model coefficients and r-squared scores
print('Coefficient:', lasso_model.coef_)
print('Scores:', lasso_model.score(X_train, y_train), lasso_model.score(X_test, y_test))

```

```
Coefficient: [ -184.8503518 16919.65968894 4601.01154766 773.00000906
 242.8726669 47.86135961 12663.89886906 -146.81819531
 9517.1057251 38.50881059 441.93502156 -261.3806402 ]
Scores: 0.7850997031385983 0.8452896114591302
```

```
In [39]: y_prediction = lasso_model.predict(X_test)
```

```
In [40]: RMSE = sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
print(RMSE)
```

```
23671.36807136452
```

```
In [41]: # K-fold CV
score_lasso = cross_val_score(lasso_model, X_train, y_train, cv=10)
score_lasso
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 159442389699.83508, tolerance: 349942370.4286453
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 240491473312.28214, tolerance: 358098534.9982602
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 257471428599.45276, tolerance: 352103963.30287695
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 234137496056.25336, tolerance: 350488682.0942842
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 224939046827.40213, tolerance: 354871608.4190128
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 235702317972.3842, tolerance: 353892070.0463491
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 228559442578.34668, tolerance: 358505522.52503693
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 234993704730.1338, tolerance: 354596261.958572
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 204457064914.1339, tolerance: 342271461.336998
```

```
model = cd_fast.enet_coordinate_descent(
```

```
C:\Users\sjue\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:53
0: ConvergenceWarning: Objective did not converge. You might want to increase the number
of iterations. Duality gap: 228607664943.06372, tolerance: 349753857.90445185
```

```
model = cd_fast.enet_coordinate_descent(
```

```
Out[41]: array([0.21617453, 0.81793705, 0.814242 , 0.80530174, 0.84199673,
 0.77720259, 0.83293262, 0.86101777, 0.84267756, 0.83439144])
```

```
In [42]: print("CV mean score:", score_lasso.mean())
print("CV mean std dev score:", score_lasso.std())
```

CV mean score: 0.7643874019363353
CV mean std dev score: 0.1840699886195586

Ridge

```
In [43]: ridge_model = Ridge().fit(X_train, y_train)
```

```
In [44]: # display model coefficients and r-squared scores
print('Coefficient:', ridge_model.coef_)
print('Scores:', ridge_model.score(X_train, y_train), ridge_model.score(X_test, y_test))
```

Coefficient: [-184.89174119 16898.96824051 4600.54796505 110.86242785
242.67866877 47.90472505 12623.20965672 -159.96286104
9500.41241026 38.55707027 441.83585603 401.52193551]
Scores: 0.7850995219712469 0.8452618243769616

```
In [45]: y_prediction = ridge_model.predict(X_test)
```

```
In [46]: RMSE = sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
print(RMSE)
```

23673.493748650915

```
In [47]: # K-fold CV
score_ridge = cross_val_score(ridge_model, X, y, cv=10)
score_ridge
```

```
Out[47]: array([0.85424213, 0.82381643, 0.84794958, 0.69904866, 0.80175182,
0.84572206, 0.86013341, 0.84208571, 0.45201179, 0.80802269])
```

```
In [48]: print("CV mean score:", score_ridge.mean())
print("CV mean std dev score:", score_ridge.std())
```

CV mean score: 0.783478428785822
CV mean std dev score: 0.11912645070598336

ElasticNet

```
In [49]: # assigning the 3 predictor variables with the highest correlation coefficient
y = df4['SalePrice']
# split data in to training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)
```

```
In [50]: elasticnet_model = ElasticNetCV().fit(X_train, y_train)
print('alpha: %f' % elasticnet_model.alpha_)
print('l1_ratio_: %f' % elasticnet_model.l1_ratio_)
```

alpha: 38192.218001
l1_ratio_: 0.500000

```
In [51]: # display model coefficients and r-squared scores
print('Coefficient:', elasticnet_model.coef_)
print('Scores:', elasticnet_model.score(X_train, y_train), elasticnet_model.score(X_test, y_test))
```

```
Coefficient: [-17.22392455  0.60378846 -0.          37.30089983 20.67096224
 74.35217833  0.          -0.          0.          6.05083074
 0.          37.29217547]
Scores: 0.4777999630212376 0.47266375696458496
```

```
In [52]: y_prediction = elasticnet_model.predict(X_test)
```

```
In [53]: RMSE = sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
print(RMSE)
```

```
43702.615614180504
```

```
In [54]: # K-fold CV
score_en = cross_val_score(elasticnet_model, X_train, y_train, cv=10)
score_en
```

```
Out[54]: array([0.19161431, 0.44163616, 0.4326376 , 0.44292176, 0.59162816,
 0.56427655, 0.56562971, 0.48415912, 0.50829106, 0.4319808 ])
```

```
In [55]: print("CV mean score:", score_en.mean())
print("CV mean std dev score:", score_en.std())
```

```
CV mean score: 0.4654775219223496
CV mean std dev score: 0.10778052509578841
```

Tuning hyperparameters for ElasticNetCV():

```
In [56]: def warn(*args, **kwargs):
pass
import warnings
warnings.warn = warn
```

```
In [57]: # define model evaluation method
cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

# define model
ratios = arange(0, 1, 0.01)

alphas = [1e-5, 1e-4, 1e-3, 1e-2, 1e-1, 0.0, 1.0, 10.0, 100.0]

model = ElasticNetCV(l1_ratio=ratios, alphas=alphas, cv=cv, n_jobs=-1);

# fit model
model.fit(X_train, y_train)
elasticnet_tuned_model = model.fit(X_train, y_train)

warnings.filterwarnings("ignore")
```

```
In [58]: # summarize chosen configuration
print('alpha: %f' % elasticnet_tuned_model.alpha_)
```

```
print('l1_ratio_: %f' % elasticnet_tuned_model.l1_ratio_)
```

```
alpha: 0.100000
l1_ratio_: 0.750000
```

```
In [59]: score_en_2 = cross_val_score(elasticnet_tuned_model, X_train, y_train, scoring='r2', cv
warnings.filterwarnings("ignore")
```

```
In [60]: print("CV mean score:", score_en_2.mean())
print("CV mean std dev score:", score_en_2.std())
```

```
CV mean score: 0.7648433160971997
CV mean std dev score: 0.18398241474890692
```

Linear Regression Models

Liner Regression Model One ([OverallQual] [GrLivArea] [GarageCars])

```
In [61]: # assigning the 3 predictor variables with the highest correlation coefficient
features = ['OverallQual', 'GrLivArea', 'GarageCars']
```

```
In [62]: X = df4[features]
```

```
In [63]: y = df4['SalePrice']
```

```
In [64]: # split data in to training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4
```

```
In [65]: # creating linear regression model
model_1 = LinearRegression().fit(X_train, y_train)
```

```
In [66]: # display model coefficients and r-squared scores
print('Coefficient:', model_1.coef_)
print('Scores:', model_1.score(X_train, y_train), model_1.score(X_test, y_test))
```

```
Coefficient: [22538.88119918    38.51336476 20167.98961928]
Scores: 0.7206411884255531 0.7741153103217534
```

```
In [67]: # predicted housing prices
y_prediction = model_1.predict(X_test)
```

```
In [68]: # model RMSE
RMSE = sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
print(RMSE)
```

```
28602.699133713417
```

```
In [69]:
```



```
# k-fold CV
scores_1 = cross_val_score(model_1, X_train, y_train, scoring='r2', cv=5)
scores_1
```

Out[69]: array([0.49854734, 0.7155773 , 0.7375631 , 0.77589921, 0.78229407])

```
In [70]: print("mean cv score:", scores_1.mean())
print("std dev cv score:", scores_1.std())
```

mean cv score: 0.7019762044171018
std dev cv score: 0.10464252297556577

Linier Regression Model Two ([OverallQual])

```
In [71]: x = df4['OverallQual'].values.reshape((-1, 1))
y = df4['SalePrice']

# split data in to training and test data
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=4)

# creating linear regression model
model_2 = LinearRegression().fit(X_train, y_train)

# display model coefficients and r-squared scores
print('Coefficient:', model_2.coef_)
print('Scores:', model_2.score(X_train, y_train), model_2.score(X_test, y_test))

# predicted housing prices
y_prediction = model_2.predict(X_test)

# model RMSE
RMSE = sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
print('RMSE:', RMSE)
```

Coefficient: [37061.76313825]
Scores: 0.5904688256569957 0.6768312441106439
RMSE: 34211.9873818558

```
In [72]: # k-fold CV
scores_2 = cross_val_score(model_2, X_train, y_train, scoring='r2', cv=5)
scores_2
```


Out[72]: array([0.45569586, 0.58430409, 0.58196803, 0.63339208, 0.65503852])

```
In [73]: print("mean cv score:", scores_2.mean())
print("std dev cv score:", scores_2.std())
```

mean cv score: 0.5820797157345505
std dev cv score: 0.06918639380526288

```
In [74]: plt.scatter(X_test, y_test, color = "blue")
plt.plot(X_train, model_2.predict(X_train), color = "red")
plt.title("OverallQual vs SalePrice")
plt.xlabel("OverallQual")
```

```
plt.ylabel("SalePrice")
plt.show()
```



Liner Regression Model Three ([qual_space])

```
In [75]: x = df4['qual_space'].values.reshape((-1, 1))
y = df4['SalePrice']

# split data in to training and test data
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=4)

# creating linear regression model
model_3 = LinearRegression().fit(X_train, y_train)

# display model coefficients and r-squared scores
print('Coefficient:', model_3.coef_)
print('Scores:', model_3.score(X_train, y_train), model_3.score(X_test, y_test))

# predicted housing prices
y_prediction = model_3.predict(X_test)

# model RMSE
RMSE = sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
print(RMSE)
```

```
Coefficient: [6.18124111]
Scores: 0.6121381394496235 0.7881827084140843
27697.739433000243
```

```
In [76]: # k-fold CV
scores_3 = cross_val_score(model_3, X_train, y_train, scoring='r2', cv=5)
scores_3
```


```
Out[76]: array([-0.48826011,  0.72447673,  0.72755186,  0.76212847,  0.73907445])
```

```
In [77]: print("mean cv score:", scores_3.mean())
print("std dev cv score:", scores_3.std())
```

```
mean cv score: 0.49299427905186183
```

std dev cv score: 0.490805507437618

```
In [78]: plt.scatter(X_test, y_test, color = "blue")
plt.plot(X_train, model_3.predict(X_train), color = "red")
plt.title("Qual_space vs SalePrice")
plt.xlabel("Qual_space")
plt.ylabel("SalePrice")
plt.show()
```



Liner Regression Model Four ([qual_space] and [GarageCars])

```
In [79]: features_2 = ['qual_space', 'GarageCars']
```

```
In [80]: x = df4[features_2]
```

```
In [81]: # split data in to training and test data
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=4)
```

```
In [82]: # creating linear regression model
model_4 = LinearRegression().fit(X_train, y_train)

# display model coefficients and r-squared scores
print('Coefficient:', model_4.coef_)
print('Scores:', model_4.score(X_train, y_train), model_4.score(X_test, y_test))
```

```
Coefficient: [4.98130032e+00 2.37770262e+04]
Scores: 0.6740395847353984 0.7992757120508557
```

```
In [83]: # predicted housing prices
y_prediction = model_4.predict(X_test)

# model RMSE
RMSE = sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
print(RMSE)
```

26962.7126297012

```
In [84]: # k-fold CV
scores_4 = cross_val_score(model_4, X_train, y_train, scoring='r2', cv=5)
scores_4
```

```
Out[84]: array([-0.24877286,  0.76434323,  0.73928354,  0.77900656,  0.77871593])
```

```
In [85]: print("mean cv score:", scores_4.mean())
print("std dev cv score:", scores_4.std())
```

```
mean cv score: 0.5625152805883451
std dev cv score: 0.405901757942075
```

When comparing the linear regression models above, the final model using [OverallQual] and [GarageCars] produced the lowest Root Mean Square Error (RMSE). However, the results appear to have slight instability. Using k-fold cross validation, Model 1 has the highest mean score and lowest standard deviation.

Testing

```
In [86]: # create dataframe using test data from kaggle
df_test = pd.read_csv("test.csv")
```

```
In [87]: # replace NaN values with zero for the test data
df_test = df_test.fillna(0)
```

```
In [88]: # sum 1st floor, 2nd floor, and basement square footage to get total square footage
sum_column = df_test['1stFlrSF'] + df_test['2ndFlrSF'] + df_test['TotalBsmtSF']

# multiply total square footage by overall quality to generate new predictor variable q
mult_column = sum_column*df_test['OverallQual']

#Sum the baths and the halfbaths
bath_column = df_test['FullBath'] + 0.5*df_test['HalfBath']

#Create a newness score (YearBuilt + (YearRemodAdd/YearBuilt))
newness_column = df_test['YearBuilt'] + (df_test['YearRemodAdd'] / df_test['YearBuilt'])

# add new predictor variables to dataframe
df_test['tot_sq'] = sum_column
df_test['qual_space'] = mult_column
df_test['baths'] = bath_column
df_test['newness_value'] = newness_column
```

Linear Regression Model One

```
In [89]: features_test = ['OverallQual', 'GrLivArea', 'GarageCars']
```

```
In [90]: X = df_test[features_test]
```

```
In [91]: test_prediction_linear = model_1.predict(X)
```

Lasso, Ridge, and ElasticNet

```
In [92]: X = df_test[name_list]
```

```
In [93]: test_prediction_lasso = lasso_model.predict(X)
test_prediction_ridge = ridge_model.predict(X)
test_prediction_elasticnet = elasticnet_tuned_model.predict(X)
```

```
In [95]: #from google.colab import files

df_linear = pd.DataFrame({'id':df_test['Id'], 'SalePrice':test_prediction_linear})
df_lasso = pd.DataFrame({'id':df_test['Id'], 'SalePrice':test_prediction_lasso})
df_ridge = pd.DataFrame({'id':df_test['Id'], 'SalePrice':test_prediction_ridge})
df_elasticnet = pd.DataFrame({'id':df_test['Id'], 'SalePrice':test_prediction_elasticne

df_linear.to_csv('group_5_linear_m3.csv', index=False)
df_lasso.to_csv('group_5_lasso_m3.csv', index=False)
df_ridge.to_csv('group_5_ridge_m3.csv', index=False)
df_elasticnet.to_csv('group_5_elasticnet_m3.csv', index=False)

# files.download('group_5_linear_m3.csv')
# files.download('group_5_lasso_m3.csv')
# files.download('group_5_ridge_m3.csv')
# files.download('group_5_elasticnet_m3.csv')
```